

Data Transfer Kit Summary

Stuart R. Slattery
Engineering Physics Department
University of Wisconsin - Madison

January 15, 2013





- Overview
- Concepts and Geometric Rendezvous
- DTK Algorithms
- Code Example

- Collection of geometry-based data mapping algorithms for shared domain problems
- Data maps allow for efficient movement of data in parallel (e.g. between meshes of a different parallel decomposition)
- Ideally maps are generated at a desirable time complexity (logarithmic)
- Input mesh and geometry data drive the map generation
- Should be viewed as a service providing suite of concrete algorithm implementations



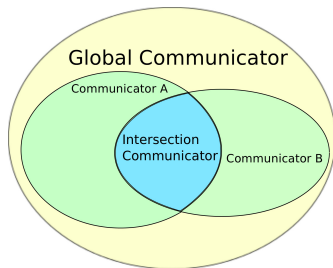
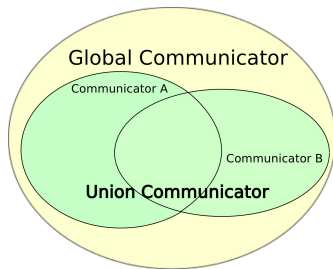
- Does not provide a general interface for all physics codes to couple to all other physics codes
- Does not provide discretization services (e.g. basis functions)
- Does not provide algorithm implementations for interface-based data transfer



- Preliminary development of mesh-based capabilities during summer 2012 CASL internship at ORNL
- Additional development of geometry-based capabilities during fall 2012
- Implemented in C++
- Heavy use of the Trilinos scientific computing libraries
- Continuous and nightly testing as part of the CASL CDash system
- Open-source BSD 3-clause license
- <https://github.com/CNERG/DataTransferKit>



- DTK handles source and target communicators of arbitrary relation
- Any amount of overlap or lack thereof supported
- A global communicator required (doesn't have to be `MPI_COMM_WORLD`)



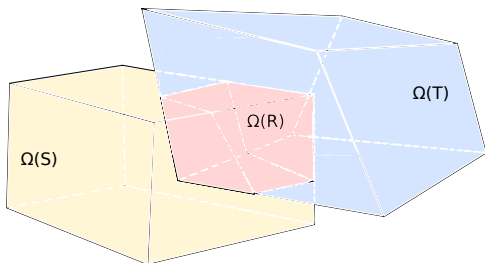


Figure: *Shared domain example.*

$\Omega(S)$ (yellow) is the source geometry,
 $\Omega(T)$ (blue) is the target geometry,
and $\Omega(R)$ (red) is the shared
domain.

- Defined over a communicator that encapsulates the union of the source and target communicators
- Source and target must be of same geometric dimension
- The rendezvous algorithm leveraged to provide parallel topology maps for shared domains

- An operator, \mathbf{M} , that defines the translation of a field, $\mathbf{F}(s)$, from a source spatial domain, Ω_S , to a field, $\mathbf{G}(t)$, in the target spatial domain Ω_T , such that $\mathbf{G}(t) \leftarrow \mathbf{M}(\mathbf{F}(s))$ and $\mathbf{M} : \mathbb{R}^D \rightarrow \mathbb{R}^D, \forall r \in \Omega_R$, where Ω_R is the geometric rendezvous of the source and target.
- \mathbf{M} is in general expensive to generate but cheap to apply
- For static Ω_S and Ω_T , building \mathbf{M} is a one-time, upfront cost

- Initially developed by the SIERRA team in the mid-2000's for parallel mesh-based data transfer ¹
- Creates a parallel topology map that can be used repeatedly for data transfer
- Map execution uses asynchronous strategy (posts and waits) with minimal messages
- Effectively $N * \log(N)$ time complexity for parallel topology map generation
- Relies on the generation of a secondary decomposition of the source and target meshes with a geometric-based partitioning (RCB)

¹S. Plimpton, B. Hendrickson, and J. Stewart, A parallel rendezvous algorithm for interpolation between multiple grids, Journal of Parallel and Distributed Computing, vol. 64, pp. 266276, 2004

The Rendezvous Decomposition

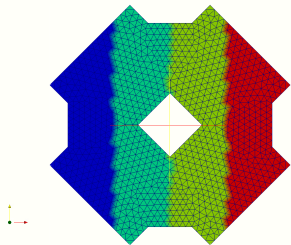


Figure: *Source mesh for 2D shared domain example.*

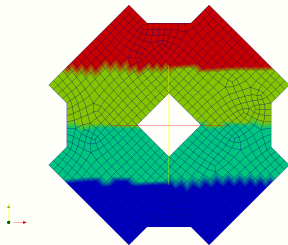


Figure: *Target mesh for 2D shared domain example.*

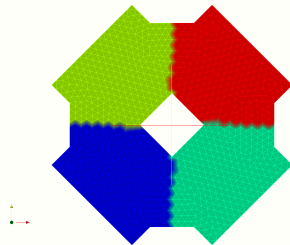


Figure: *Rendezvous decomposition for 2D shared domain example.*



- Hierarchical parallel search tree
- Rendezvous decomposition provides parallel search
- kD-tree provides on-process proximity search
- Newton iterations provide final point location
- Results in reasonable scalability

- Mesh-to-mesh transfer
- Worst case scenario study (all-to-all) with random points
- Qualitatively similar to the SIERRA results
- Largest test problems so far over $1.0E9$ elements and $1.0E5$ cores




Figure: *Local mesh partition for scaling studies. This particular mesh partition has $1.0E4$ tri-linear hexahedrons.*

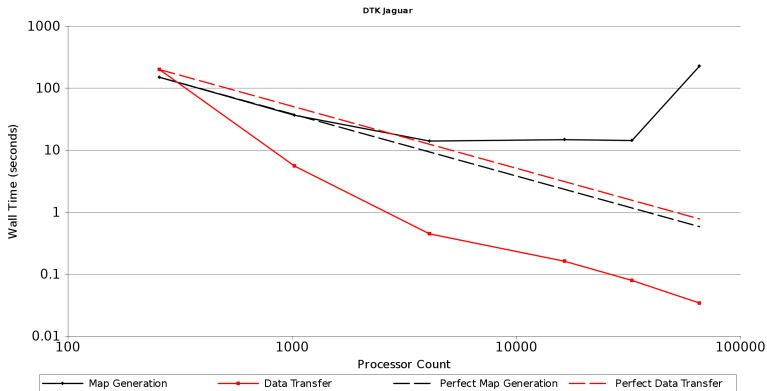


Figure: Strong scaling study results. The solid black curve reports the wall time to generate the mapping vs. number of processors while the solid red curve reports the wall time to transfer the data vs. number of processors. The dashed lines give perfect strong scaling the map generation (black) and the data transfer (red).

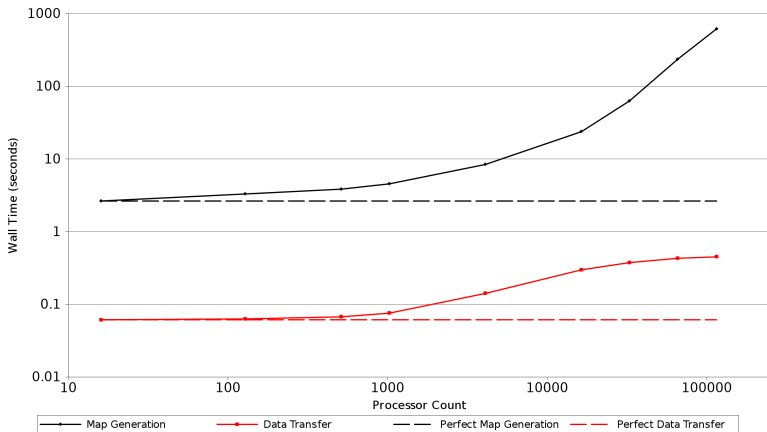


Figure: Weak scaling study results. The solid black curve reports the wall time to generate the mapping vs. number of processors while the solid red curve reports the wall time to transfer the data vs. number of processors. The dashed lines give perfect weak scaling the map generation (black) and the data transfer (red).



- Meshes are viewed as geometric structures
- A subset of total mesh information is needed:
 - Vertex coordinates
 - Element topology
 - Element connectivity
 - Connectivity permutation
- Parallel information is not required
- A communicator and global IDs are required

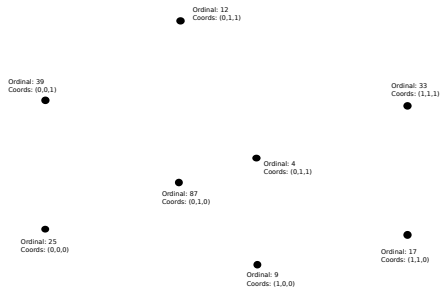


Figure: Basic vertex description for a mesh.

Each vertex is required to have a unique global ID

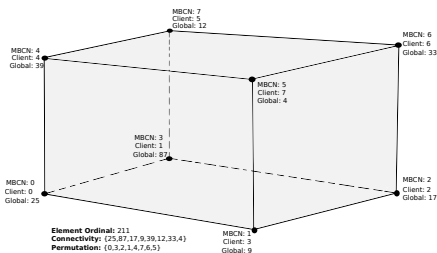



Figure: Basic element description for a mesh.

Each element is required to have a unique global ID


Connectivity Permutation




LINE SEGMENT




TRIANGLE




QUADRILATERAL




TETRAHEDRON



PYRAMID



WEDGE



HEXAHEDRON

- Allow for user specification of canonical ordering
- Permutation list specifies the variation in ordering for a specified topology
- Other mesh topologies currently not supported

Figure: Canonical vertex connectivity schemes for elements in DTK.

- All topologies in a mesh must be of the same dimension
- Each topology contained in a block
- Can have multiple blocks of the same topology

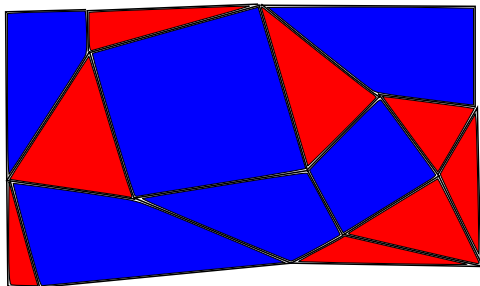


Figure: Hybrid mesh example.

Quadrilaterals (blue) must be specified in a different mesh block than the triangles (red). Both blocks can contain the mutual mesh vertices that construct their elements.



- DTK's perception of general geometric structures is primitive
- A more capable geometry engine required for advanced algorithms
- A subset of total geometry information is needed:
 - Centroid
 - Bounding box
 - Measure
 - Point inclusion
- Parallel information is not required
- A communicator and global IDs are required

- Actual discretization of the field is not explicitly formulated
- Access to discretization of fields is generated through user code function evaluations at points in physical space:

$$\hat{f} \leftarrow \mathbf{F}(\hat{r}), \forall \hat{r} \in \Omega$$

- In the context of Ω discretized by a mesh, these evaluations can instead be written in terms of a single mesh element, $\omega \in \Omega$:

$$\hat{f} \leftarrow \mathbf{F}(\hat{r}), \forall \hat{r} \in \omega$$



- What user code should expect:
 - A list of global element IDs that are on-process (ω)
 - A corresponding list of point coordinates (\hat{r})
- What user code should provide:
 - A list of the resulting function evaluations for each global element ID/point coordinate pair provided (\hat{f})
- All data provided to user code will be local with respect to input data
- C++ inheritance (mix-in interface)

- Consider a measure-weighted integral:

$$f_{\Omega} = \frac{\int_{\Omega} \mathbf{F}(r) dr}{\int_{\Omega} dr}$$

- In the context of Ω discretized by a mesh, these evaluations can instead be written in terms of a single mesh element, $\omega \in \Omega$:

$$f_{\omega} = \int_{\omega} \mathbf{F}(r) dr$$

- The integral over Ω will be the measure-weighted summation of all element integrals:

$$f_{\Omega} = \frac{1}{m_{\Omega}} \sum_i f_{\omega_i}, \quad \forall \omega_i \in \Omega$$



- Element-wise spatial integrals generated through user code




- What user code should expect:
 - A list of global element IDs that are on-process (ω)
- What user code should provide:
 - A list of the resulting function integrations for each global element ID provided (f_ω)
- All data provided to user code will be local with respect to input data
- C++ inheritance (mix-in interface)



- What user code should provide:
 - Allocated memory block large enough to hold the resulting function evaluations or integrations for the specified local target objects
- What user code should expect:
 - No allocation of memory
 - A result of 0 if no evaluation or integration occurred for that object
- All data provided to user code will be local with respect to input data
- Implemented as array views (pointer and size)




- Mesh-to-Mesh transfer ^a
- Used to move $\mathbf{F}(\hat{r})$ between meshes of arbitrary distribution
- Requires user code for evaluations in mesh elements

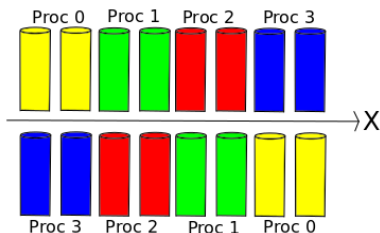


^aExample provided by Roger Pawlowski

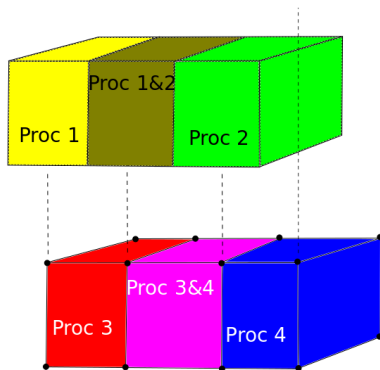
- Mesh-to-geometry transfer
- Used to assemble f_{Ω} with mesh and geometry of arbitrary distribution into measure-weighted integral
- The mesh is assumed conformal
- Requires user code for integrations in mesh elements
- See example/IntegralAssembly



Other Rendezvous-Based Maps: Geometry to Geometry




- Simple geometry-to-geometry transfer capability
- Geometries are assumed conformal
- Requires user code for evaluations in geometry
- See `example/GeometryToGeometry`



- Similar to mesh-based rendezvous
- Does not require a mesh, conceptual in this case
- Requires user code for evaluations in geometry
- See [example/GeometryToMesh](#)

- examples/WaveDamper
- Using DTK in the context of a Picard iteration
- Start with existing codes, expose data to DTK
- Use mesh-to-mesh mapping in 1D



Super Simple Example: Wave



- 1D code
- Initial conditions: $\mathbf{F}(x) = \cos(x)$

```
class Wave {
private:
    Teuchos::RCP<const Teuchos::Comm<int> > comm;
    Teuchos::RCP<std::vector<double> > grid;
    Teuchos::RCP<std::vector<double> > data;
    Teuchos::RCP<std::vector<double> > damping;

public:
    Wave( Teuchos::RCP<const Teuchos::Comm<int> > _comm,
          double x_min, double x_max, int num_x)
    {
        // Create the grid.
        // Set initial conditions.
    }

    void solve()
    { /* Apply the damping to the local data */ }
};
```

Super Simple Example: Damper



- 1D code
- Initial conditions: none

```
class Damper
{
    private:
        Teuchos::RCP<const Teuchos::Comm<int>> > comm;
        Teuchos::RCP<std::vector<double>> > data;
        Teuchos::RCP<std::vector<double>> > damping;
        Teuchos::RCP<std::vector<double>> > grid;

    public:
        Damper( Teuchos::RCP<const Teuchos::Comm<int>> > _comm,
                double x_min, double x_max, int num_x)
        { /* Create the grid. */ }

        void solve()
        { /* Apply damping to the local data. */ }
};
```


- Functions added to the Wave class

```
// Get the communicator.
 Teuchos::RCP<const Teuchos::Comm<int> > get_comm() const
{ return comm; }

// Get a const reference to the local grid.
const Teuchos::RCP<std::vector<double> > get_grid() const
{ return grid; }

// Get a reference to the local data.
const Teuchos::RCP<std::vector<double> > get_data() const
{ return data; }

// Get a reference to the local data space storing
// the damping coefficients.
Teuchos::RCP<std::vector<double> > get_damping()
{ return damping; }
```

- Functions added to the Damper class

```
// Get the communicator.
 Teuchos::RCP<const Teuchos::Comm<int> > get_comm() const
{ return comm; }

// Get a reference to the local damping data.
 Teuchos::RCP<std::vector<double> > get_damping() const
{ return damping; }

// Get a reference to the local grid.
 Teuchos::RCP<std::vector<double> > get_grid() const
{ return grid; }

// Get a reference to the memory space for external data
// to be applied to.
 Teuchos::RCP<std::vector<double> > get_external_data()
{ return data; }
```

Super Simple Example: Wave Function Evaluations



- Inherit from `DataTransferKit::FieldEvaluator`
- Implement the `Evaluate()` function

```
class WaveEvaluator :  
public DataTransferKit::FieldEvaluator<ordinal_type, field_type>  
{  
    public:  
    WaveEvaluator( const RCP_Wave& wave );  
  
    field_type evaluate(const Teuchos::ArrayRCP<int>& elements,  
                        const Teuchos::ArrayRCP<double>& coords)  
    {  
        // Setup an output field.  
        // Get the Wave grid.  
        // Get the Wave data.  
  
        // Interpolate the Wave data onto the given  
        // coordinates using a linear basis. The  
        // coordinates will be valid for the given elements.  
    }  
};
```

Super Simple Example: Damper Function Evaluations

- Inherit from DataTransferKit::FieldEvaluator
- Implement the Evaluate() function

```
class DamperEvaluator :
public DataTransferKit::FieldEvaluator<ordinal_type, field_type>
{
    public:
        DamperEvaluator( const RCP_Damper& damper );

        field_type evaluate(const Teuchos::ArrayRCP<int>& elements,
                           const Teuchos::ArrayRCP<double>& coords)
        {
            // Setup an output field.
            // Get the Damper grid.
            // Get the Damper data.

            // Interpolate the Damper data onto the given
            // coordinates using a linear basis. The
            // coordinates will be valid for the given elements.
        }
};
```

- Provide DTK mesh and field data

```
class WaveAdapter
{
    // Get the wave mesh.
    static RCP<DataTransferKit::MeshManager<MeshType> >
    getMesh( const RCP_Wave& wave );

    // Get the wave field evaluator.
    static RCP_Evaluator
    getFieldEvaluator( const RCP_Wave& wave );

    // Get the wave target coordinates from the mesh.
    static RCP<DataTransferKit::FieldManager<MeshType> >
    getTargetCoords( const RCP_Wave& wave );

    // Get the wave target space.
    static RCP<DataTransferKit::FieldManager<FieldType> >
    getTargetSpace( const RCP_Wave& wave );
};
```

- Provide DTK mesh and field data

```
class DamperAdapter
{
    // Get the damper mesh.
    static RCP<DataTransferKit::MeshManager<MeshType> >
    getMesh( const RCP_Damper& damper );

    // Get the damper field evaluator.
    static RCP_Evaluator
    getFieldEvaluator( const RCP_Damper& damper );

    // Get the damper target coordinates from the mesh.
    static RCP<DataTransferKit::FieldManager<MeshType> >
    getTargetCoords( const RCP_Damper& damper );

    // Get the damper target space.
    static RCP<DataTransferKit::FieldManager<FieldType> >
    getTargetSpace( const RCP_Damper& damper );
};
```

Super Simple Example: Driver Setup



```
int main(int argc, char* argv[])
{
    // Parallel setup.
    Teuchos::RCP<const Teuchos::Comm<int> > comm_union;
    DataTransferKit::CommTools::unite(
        wave_comm, damper_comm, comm_union );
    // Wave setup.
    // Damper setup.
    // Mapping.
    DataTransferKit::SharedDomainMap<WaveAdapter::MeshType,
                                    DamperAdapter::MeshType>
        wave_to_damper_map( comm_union, 1 );
    wave_to_damper_map.setup(wave_mesh, damper_target_coords);

    DataTransferKit::SharedDomainMap<DamperAdapter::MeshType,
                                    WaveAdapter::MeshType>
        damper_to_wave_map( comm_union, 1 );
    damper_to_wave_map.setup(damper_mesh, wave_target_coords);

    // Solve. (next slide)
}
```



- Picard iterations to convergence

```
while( norm > tolerance && num_iter < max_iter )
{
    wave_to_damper_map.apply( wave_evaluator ,
                              damper_target_space );

    damper->solve();

    damper_to_wave_map.apply( damper_evaluator ,
                              wave_target_space );

    wave->solve();

    // Compute Wave data norm.
    // Update the iteration count.
    // Barrier before proceeding to the next iteration.
}
```